

AAHAN JAIN

Seattle, WA | +1 (919) 650-7852 | aahanjain1114@gmail.com | [linkedin/aahan-jain](https://www.linkedin.com/in/aahan-jain) | [github/aahan14jain](https://github.com/aahan14jain) | [portfolio](#)

TECHNICAL SKILLS

Languages: Python, Java, C++, TypeScript, SQL

AI/ML & GenAI: Agentic RAG, LangGraph, LangChain, LlamaIndex, OpenAI API, Azure OpenAI, AWS Bedrock, Pinecone

Machine Learning & Statistics: scikit-learn, PyTorch, XGBoost, Prophet, SHAP, statsmodels, scipy, TensorFlow

Full-Stack & Data: PostgreSQL, MongoDB, Redis, DynamoDB, RedShift, FastAPI, Django, Next.js, React.js

Cloud, DevOps & Tools: AWS (ECS Fargate, Lambda, S3, CloudWatch), Terraform, Docker, Jenkins, CI/CD, Git, Cursor, GitHub Copilot, Airtable

EXPERIENCE

Research — AI-Accelerated Software Engineering (DocuMind) | [GitHub](#)

Sept 2025 – Feb 2026

George Mason University, CS 691

Fairfax, VA

- Built DocuMind, a **LangGraph** agentic RAG system indexing multi-language repos into vector and graph representations, replacing manual code exploration with natural-language Q&A grounded in file:line citations.
- Engineered a hybrid retrieval pipeline **OpenAI** embeddings in **Pinecone**, **LlamaIndex** BM25 fusion/re-ranking, **NetworkX** graph traversal serving code Q&A at **2.8s** average latency.
- Benchmarked 3 retrieval strategies via **RAGAS**, improving answer faithfulness **14%** (**0.64** to **0.73**), validated by **207+** automated tests across **12** API endpoints.

Software Development Intern

Aug 2023 – April 2024

Hexagon, India

Hyderabad, India

- Shipped a production AI customer-support assistant on prompt-engineered **RAG pipelines** and **Azure OpenAI** to clear a growing backlog of repetitive tickets, fielding 500+ daily queries at 90% response accuracy.
- Orchestrated **LangChain** LLM workflows behind **CI/CD** automation on **Azure** to attack a slow 4-day resolution cycle, cutting resolution time to 2 days and stabilizing releases across 10+ client environments.
- Architected backend services in **Django** and **MongoDB** to remove latency bottlenecks in real-time AI query handling, lifting API response performance by 5%.

PROJECTS

86IQ — AI Decision Platform for Restaurant Owners | [GitHub](#)

Mar 2026 – Present

- Designed a multi-agent **LangGraph** pipeline (context, decision, explanation agents) behind a **FastAPI** backend, serving prep/pricing/menu actions through a **Next.js** dashboard.
- Diagnosed data leakage inflating an **XGBoost/PyTorch** ensemble to 99% accuracy, corrected to a realistic **74.5%** `f1_macro`, then shipped **Prophet** (**2.6x** better than baseline).
- Delivered **SHAP**-explained (32ms) predictions with **AWS Bedrock**-polished summaries, deployed on Terraform-provisioned **AWS** (ECS Fargate, Lambda-triggered retraining).

Cadence — Customer Recommendation Engine | [GitHub](#)

Jan 2025 – May 2025

- Designed a full-stack app (**React**, **Django** REST API, **PostgreSQL**) replacing generic, mistimed promotions with behavior-driven, vendor-matched offers, secured via **JWT** auth.
- Trained a Machine Learning model (Random Forest) achieving **90%** accuracy predicting repurchase timing (soon/medium/late) across **3.3K** labeled purchase gaps, ranking each customer's next offer.
- Served confidence-scored recommendations via a **JWT**-scoped **Django** API with vendor-matching and a **DynamoDB** 1-hour cache, syncing results to **Airtable** for review.

Lift — A/B Testing & LLM Evaluation Framework | [GitHub](#)

Jun 2025 - Sept 2025

- Built a **Python** A/B framework simulating 10K synthetic users (**pandas**) across control + 5 variants, running z-tests and Welch's t-tests (**scipy**, **statsmodels**) with FDR correction.
- Modeled conversion likelihood via **scikit-learn** logistic regression, sourcing production experiment data from **Redshift** and persisting results to **S3** for reproducibility.
- Extended to LLM evaluation via paired t-tests across **5,200** inputs x 5 prompt variants, with an **S3**-triggered **Lambda** re-running the pipeline and **CloudWatch** monitoring.

EDUCATION

George Mason University

Fairfax, VA

Master of Science in Computer Science, GPA: 3.73/4.0

Aug 2024 – May 2026

Vishwakarma Institute of Technology

Pune, India

Bachelor of Technology in Computer Engineering, GPA: 3.37/4.0

Aug 2020 – May 2024