

AAHAN JAIN

Seattle, WA | +1 (919) 650-7852 | aahanjain1114@gmail.com | [linkedin/aahan-jain](https://www.linkedin.com/in/aahan-jain) | [github/aahan14jain](https://github.com/aahan14jain) | [portfolio](#)

TECHNICAL SKILLS

Languages: Python, Java, C++, TypeScript, SQL, JavaScript, HTML/CSS

Frontend & Backend: React.js, Next.js, Angular, FastAPI, Django, Spring Boot, REST APIs, Apache Kafka

Databases: PostgreSQL, MongoDB, Redis, DynamoDB

Testing: JUnit, Mockito

Cloud & DevOps: AWS (EC2, S3, CloudWatch, Lambda), Azure, Terraform, Docker, Kubernetes, CI/CD, Git, Linux, Splunk, Prometheus, Grafana

AI/ML Engineering & Tooling: Agentic RAG, LangGraph, LangChain, LlamaIndex, OpenAI API, Azure OpenAI, AWS Bedrock, Pinecone, Cursor, GitHub Copilot

EXPERIENCE

Research — AI-Accelerated Software Engineering (DocuMind) | [GitHub](#)

Sept 2025 – Dec 2025

George Mason University, CS 691

Fairfax, VA

- Engineered DocuMind, a **Python/Flask** + **LangGraph** agentic RAG service built with **Cursor** and **GitHub Copilot**, cutting repo onboarding from hours of manual tracing to seconds of natural-language queries.
- Combined **OpenAI** embeddings, **Pinecone** vector storage, and **LlamaIndex** BM25 fusion/re-ranking into a hybrid retrieval layer serving code Q&A at **2.8s** average query latency.
- Persisted chat history in **MongoDB** and generated artifacts in **S3**, containerized with **Docker**, backed by **207+** automated tests across **12** API endpoints.

Software Development Intern

Aug 2023 – April 2024

Hexagon, India

Hyderabad, India

- Shipped a production AI customer-support assistant on prompt-engineered **RAG pipelines** and **Azure OpenAI** to clear a growing backlog of repetitive tickets, fielding 500+ daily queries at 90% response accuracy.
- Orchestrated **LangChain** LLM workflows behind **CI/CD** automation on **Azure** to attack a slow 4-day resolution cycle, cutting resolution time to 2 days and stabilizing releases across 10+ client environments.
- Architected backend services in **Django** and **MongoDB** to remove latency bottlenecks in real-time AI query handling, lifting API response performance by 5%.

PROJECTS

Choreo — Event-Driven Checkout Microservices | [GitHub](#)

Jan 2026 – April 2026

- Engineered checkout into **Spring Boot** microservices (order, inventory, payment) coordinated by a choreographed **Saga** over **Kafka**, keeping transactions consistent across 1,000+ concurrent orders.
- Eliminated duplicate event processing under load with **Redis SETNX+TTL** idempotency running on **Kubernetes**, validated with **JUnit/Mockito** suites at 85% coverage before release.
- Deployed via **GitHub Actions** to **Kubernetes**, instrumenting the system with **CloudWatch** (IAM-scoped) and **Prometheus/Grafana** to trace a latency regression to source, cutting read latency 35%.

86IQ — AI Decision Platform for Restaurant Owners | [GitHub](#)

Mar 2026 – Present

- Architected a multi-agent **LangGraph** pipeline (context, decision, and explanation agents) behind a **FastAPI** backend, serving ranked daily prep/pricing/menu actions through a **Next.js** dashboard.
- Caught data leakage inflating model accuracy to **99%**, corrected it to a realistic **75%**, then shipped the best-performing forecaster (**2.6x** lower error than baseline) with **AWS Bedrock**-generated explanations.
- Deployed to **AWS** via **Terraform** (**ECS Fargate**, **Lambda**-triggered retraining) with **GitHub Actions** CI/CD and **40+** automated tests validating every release.

Cadence — Customer Recommendation Engine | [GitHub](#)

Jan 2025 – May 2025

- Designed a full-stack app (**React**, **Django** REST API, **PostgreSQL**) replacing generic, mistimed promotions with behavior-driven, vendor-matched offers, secured via **JWT** auth.
- Trained a Machine Learning model (Random Forest) achieving **90%** accuracy predicting repurchase timing (soon/medium/later) across **3.3K** labeled purchase gaps, ranking each customer's next offer.
- Built **JWT**-scoped recommendation APIs over **22.5K** transactions with a **DynamoDB** 1-hour cache, packaged via **Docker** with **Splunk** logging and **Jenkins** CI.

EDUCATION

George Mason University

Fairfax, VA

Master of Science in Computer Science, GPA: 3.73/4.0

Aug 2024 – May 2026

Vishwakarma Institute of Technology

Pune, India

Bachelor of Technology in Computer Engineering, GPA: 3.37/4.0

Aug 2020 – May 2024